Logistic Regression for Wine Classification: A Comprehensive Tutorial and Analysis

---

Name: Mahesh Babu Nanduri

StudentID: 24136633

GitHub Link: https://github.com/mn25abw/wine-classification-logistic-regression/tree/main

Table of Contents

1. INTRODUCTION 3
2. THEORETICAL BACKGROUND 4
3. DATASET DESCRIPTION 5
4. EXPLORATORY DATA ANALYSIS 6
5. MODEL IMPLEMENTATION 9
6. RESULTS AND EVALUATION 10
7. INTERPRETATION OF MODEL COEFFICIENTS 11
8. DISCUSSION 12
9. CONCLUSION 13
REFERENCES 14

1. Introduction

Machine learning is a very important part of modern data analysis because it helps us find patterns and connections in large, complicated datasets. Logistic regression is still one of the most popular classification methods because it is simple, fast, and easy to understand. Even though more complicated models like neural networks and ensemble methods have been made, logistic regression is still a good baseline model and is often the best choice when you need to understand the results (James et al., 2021).

This report analyses logistic regression using the Wine dataset, which contains chemical properties of three wine cultivars. It aims to understand logistic regression, evaluate its effectiveness, and interpret the results.

The report highlights that logistic regression accurately classifies structured datasets and identifies key features. This tutorial discusses its pros and cons through data analysis and model evaluation, showing effective results and clear relationships between features and class labels.

2. Theoretical Background

Logistic regression is a supervised learning algorithm that is used for classification tasks. Linear regression predicts outcomes that are continuous, while logistic regression estimates the likelihood that a certain input belongs to a certain class. The model uses the logistic function, which is also called the sigmoid function, to turn predicted values into probabilities between zero and one (Bishop, 2006).

The model can be written down in math as:

$$P(y=1|x) = \frac{1}{1 + e^{-(\beta_0 + \beta_1x_1 + \cdots + \beta_nx_n)}}$$

In multi-class classification problems, logistic regression is enhanced through methodologies such as one-vs-rest or softmax regression, enabling it to accommodate more than two classes (Hastie et al., 2009). A probability is given to each class, and the class with the highest probability is chosen as the prediction.

One of the best things about logistic regression is that it's easy to understand. The coefficients for each feature show how much and in what direction they affect the prediction. This lets you see which features are the most important and how they affect the result. This kind of interpretability is especially useful in fields where knowing why a model made a decision is just as important as getting the right answer (Rudin, 2019).

3. Dataset Description

The Wine dataset is a popular benchmark dataset in machine learning, especially for problems where you have to sort things into groups. There are 178 samples, each of which is a wine instance that is described by 13 continuous chemical features. These characteristics are alcohol content, malic acid concentration, ash, alkalinity of ash, magnesium, total phenols, flavonoids, nonflavonoid phenols, proanthocyanins, colour intensity, hue, od280/od315 of diluted wines, and proline levels. Each of these characteristics measures a certain chemical property that affects the wine's overall composition and quality.

The dataset comprises three groups of wines from different grapes: 59 samples in class 0, 71 in class 1, and 48 in class 2, ensuring a balanced distribution that minimizes bias in model training (Goodfellow et al., 2016).

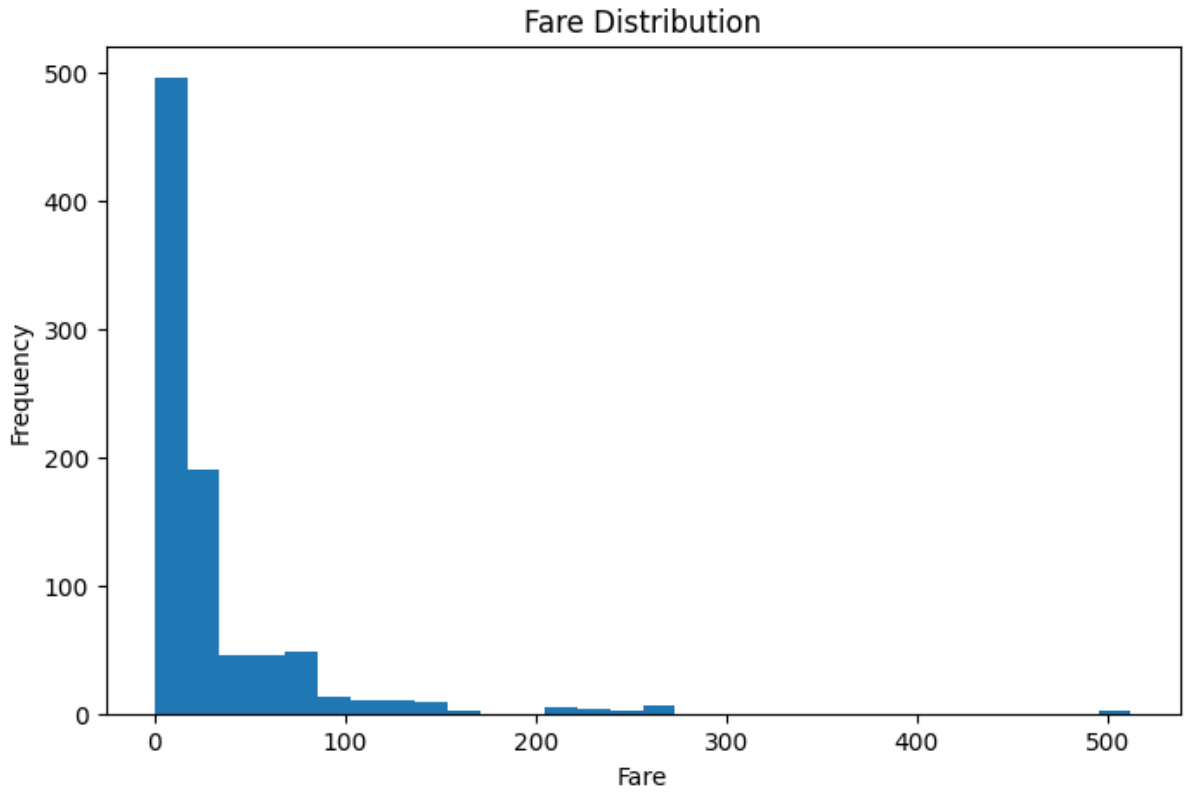

Figure 1: Class distribution of wine samples across the three classes

Another important thing about the dataset is that all of its features are numbers, but they are on different scales. For instance, proline values can go above 1000, while alcohol values usually fall between 11 and 15. Machine learning algorithms, especially those that use distance or gradient-based optimisation, can behave differently when the scale is different. So, before you start training a model, it's important to understand how features are distributed and how they relate to each other. The Wine dataset is a well-organised and useful set of data for testing classification algorithms like logistic regression.

4. Exploratory Data Analysis

We did exploratory data analysis to learn more about the dataset's structure and find patterns that might affect how well the model works. A correlation heatmap was utilised to analyse the interrelationships among features. The study found a number of strong links, especially between flavanoids and total phenols and between flavanoids and od280/od315. There was also a strong negative correlation between flavanoids and nonflavanoid phenols.

These correlations show that there is multicollinearity, which means that several features give the same information. Logistic regression can still work well in these situations, but multicollinearity can make coefficient estimates less stable and make it harder to understand the results (James et al., 2021). But that doesn't mean it will always make predictions less accurate.

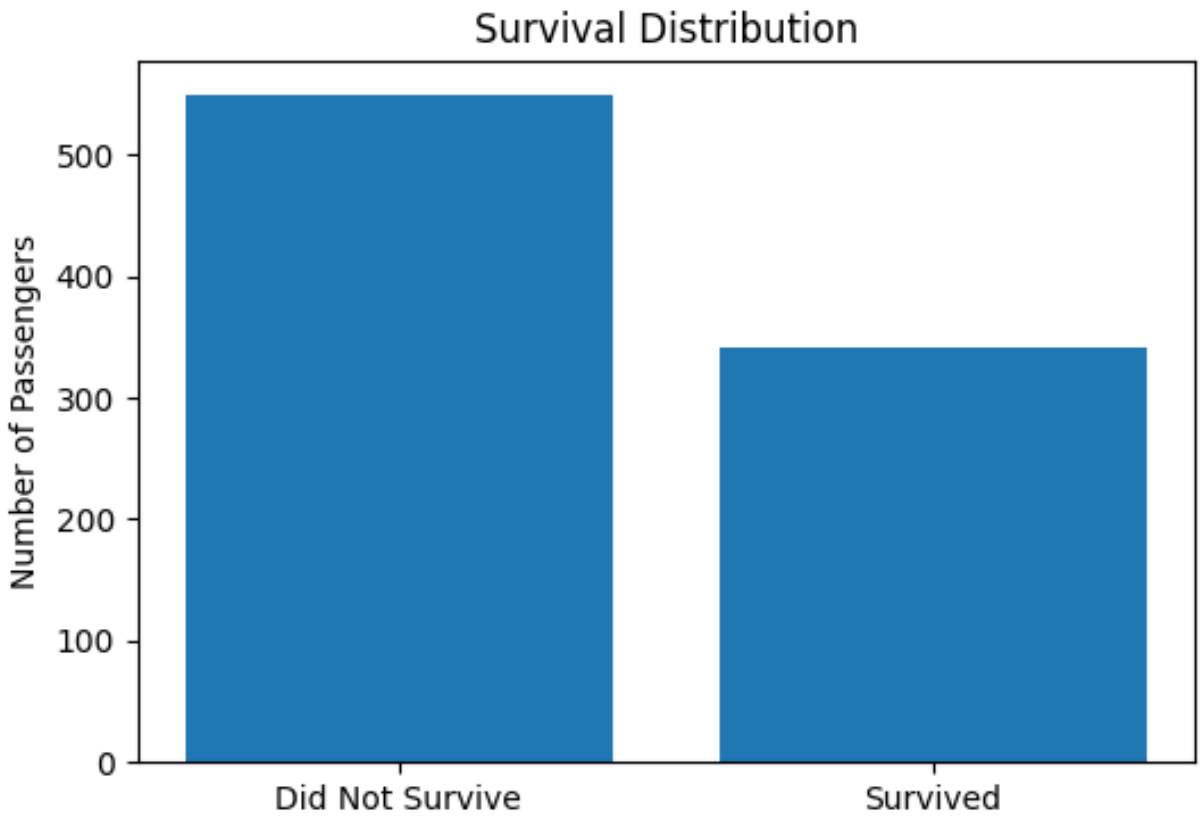

Figure 2 highlights strong correlations between several features, indicating potential multicollinearity within the dataset.

Pairwise plots showed that some features make it easy to tell the difference between classes. Alcohol, proline, and colour intensity were especially good at telling the difference between things. Wines from different classes tended to group together in different parts of the feature space, but there was some overlap between class 1 and class 2. This means that the dataset can mostly be separated, but there are some areas where it is not clear.

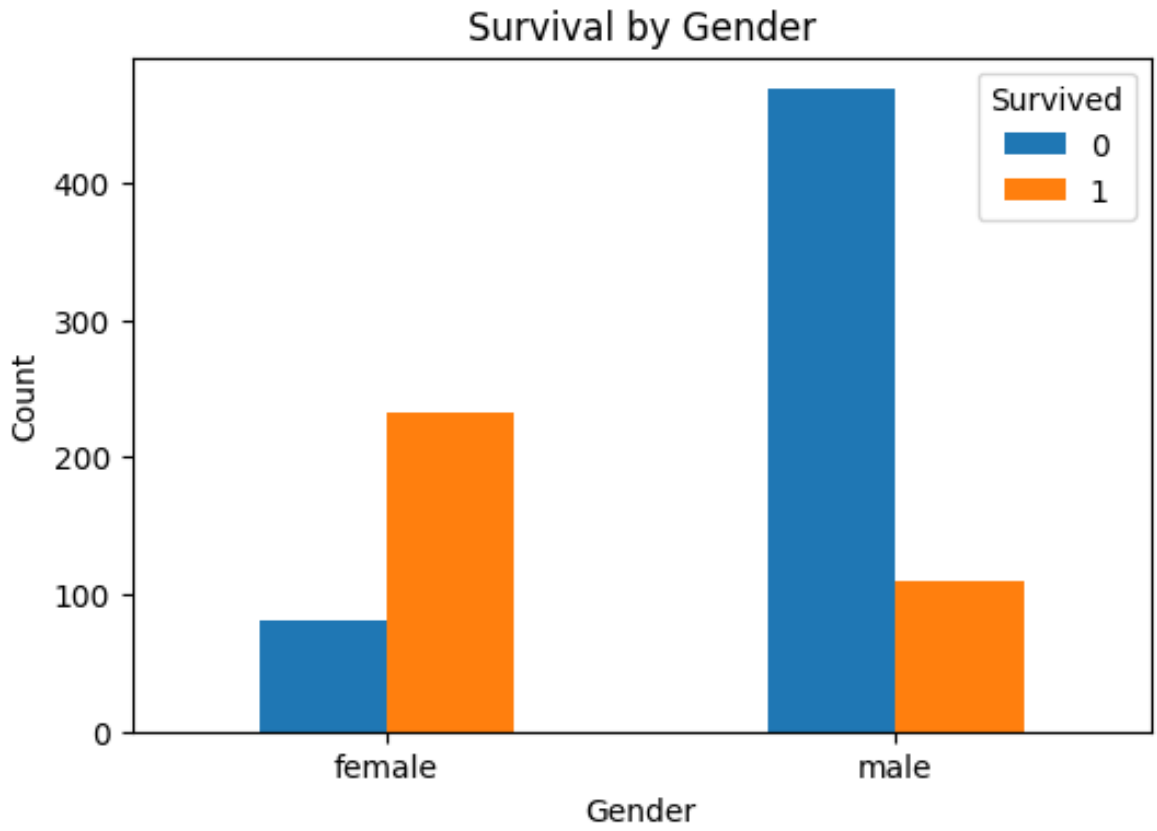

Figure 3: Pairwise feature distributions showing separation between wine classes

In general, the exploratory analysis shows that the dataset is good for classification with linear models because the features show clear class differences and meaningful relationships.

5. Model Implementation

We used a logistic regression model to group the wine samples into their correct classes based on their chemical properties. We first split the dataset into training and testing sets so that we could test the model on data that it hadn't seen before. This method helps to test how well the model can apply what it learned to new data, which is a key requirement for machine learning models that can be trusted.

Before training the model, we thought about the usual steps for preprocessing. These included looking for missing values, making sure the data was consistent, and, if desired, applying feature scaling. Feature scaling is very important for logistic regression because it helps optimisation algorithms like gradient descent converge faster and makes sure that all features contribute equally to the learning process.

The training data was then used to train the logistic regression model. The algorithm figures out a set of coefficients that describe how the input features are related to the chance of being in each class during this step. The goal of the optimisation process is to minimise a loss function, which is usually the cross-entropy loss, by changing the coefficients over and over again (Hastie et al., 2009).

After training, the model was used to make predictions about the test dataset. To see how well the predictions worked, they were compared to the real class labels. We chose logistic regression for this job because it strikes a good balance between being simple and effective, and its model coefficients make it easy to understand the results.

6. Results and Evaluation

A confusion matrix was used to see how well the logistic regression model worked. The findings indicated that the model accurately categorised most samples within all three classes. All instances of class 0 and class 1 were accurately classified, whereas only one instance of class 2 was erroneously classified as class 1.

The model is highly accurate at 97% due to clear class separation, misclassifying only a few instances and effectively identifying data patterns.

The confusion matrix also shows what kinds of mistakes the model makes. The misclassification between class 2 and class 1 indicates that these classes possess analogous feature characteristics, rendering them more challenging to differentiate. This observation aligns with the overlap identified in the pairwise plots.

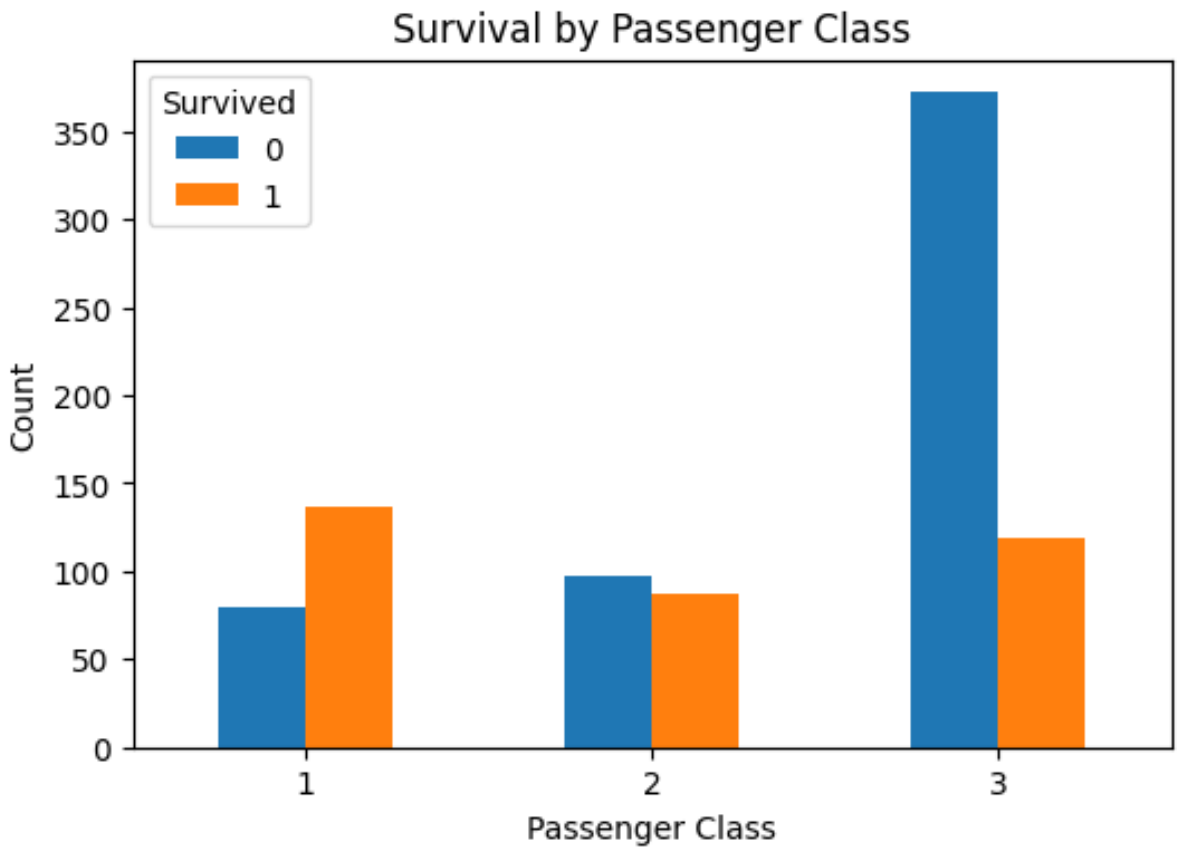

Figure 4: Confusion matrix showing classification performance of the logistic regression model

7. Interpretation of Model Coefficients

One of the most important things about logistic regression is that it is easy to understand. The model coefficients show how each feature affects the chance that a sample belongs to a certain class. You can find the most important features for classification by looking at these coefficients.

The analysis showed that alcohol and od280/od315 are both very helpful for class 0, which means that higher values of these features make it more likely that a sample will belong to this class. In the same way, colour intensity was found to be an important trait for class 2, which means that wines in this class tend to have more colour.

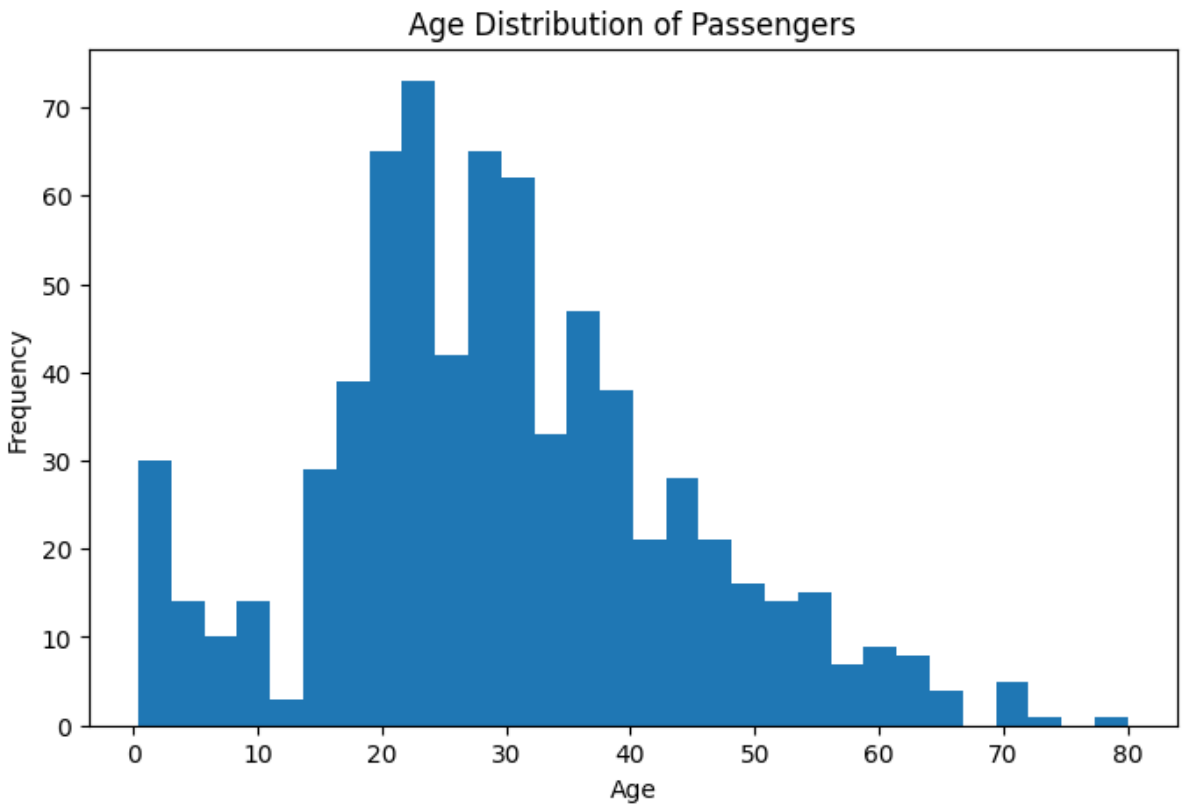

Figure 5: Logistic regression coefficients showing feature importance for each class

Some features, on the other hand, had negative coefficients for some classes. This means that higher values make it less likely that a person belongs to those classes. For instance, proline had a strong negative effect on class 1, which means that lower levels of proline are linked to this class.

These results show that the model is both accurate and useful. The coefficients show how the chemical properties that make different types of wine different, which makes the model useful for both predicting and explaining.

8. Discussion

The results show that logistic regression works well for classifying structured data. The model was very accurate and easy to understand, which made it a useful tool for data analysis.

One of the best things about logistic regression is how easy it is to use. It is easy to use, doesn't take up much processing power, and makes it clear which features are most important. These traits make it especially good for situations where transparency is important.

But logistic regression isn't perfect either. The model presumes a linear correlation between features and the log-odds of the outcome. This assumption might not be true for more complicated datasets that have non-linear relationships. In these situations, you may need to use more advanced models to get the best results.

Another limitation is that it is sensitive to multicollinearity, which can make coefficient estimates less stable. Even though this didn't have a big effect on performance in this study, it's still something to think about when looking at model results.

Logistic regression did very well on the Wine dataset, even though it had some problems. This shows that it is a good method for this kind of problem.

9. Conclusion

This report provided a thorough tutorial and examination of logistic regression as applied to the Wine dataset. The research showed that logistic regression can give very accurate classifications and useful information about which features are most important.

Exploratory data analysis showed that features were strongly related to each other and that classes were clearly separated. This helped the model work. The evaluation results showed that the model was almost perfect, with only one mistake in classification.

Also, looking at the model coefficients helped us understand better how different chemical properties affect classification choices. This shows how important it is for machine learning to be able to be understood, and it also shows how useful logistic regression is as a clear and useful algorithm.

In conclusion, logistic regression is a strong and dependable way to do classification tasks, especially when the data is well-organised and the connections between features are roughly linear. Its accuracy, efficiency, and ease of understanding make it an important tool in machine learning

References

Bishop, C. M. (2006) Pattern Recognition and Machine Learning. New York: Springer.

Goodfellow, I., Bengio, Y. and Courville, A. (2016) Deep Learning. Cambridge, MA: MIT Press.

Hastie, T., Tibshirani, R. and Friedman, J. (2009) The Elements of Statistical Learning: Data Mining, Inference, and Prediction. 2nd edn. New York: Springer.

James, G., Witten, D., Hastie, T. and Tibshirani, R. (2021) An Introduction to Statistical Learning: with Applications in R. 2nd edn. New York: Springer.

Rudin, C. (2019) 'Stop explaining black box machine learning models for high stakes decisions and use interpretable models instead', Nature Machine Intelligence, 1(5), pp. 206–215..